

# Forecasting Stock Volatility

A Model Comparison for Louis, Quantitative Trader

Presenter: Yun Chang 520397297

# Project Overview & Data Description

**Goal:** Predict realized stock volatility from high-frequency data.

**Data Used:** First 500 time\_id sessions from stock\_1.csv

## Key Features Created:

- **WAP (Weighted Average Price):** Combines bid/ask prices & sizes.
- **Bid-Ask Spread:** Relative measure of liquidity.
- **Log Returns:** Calculated per second from WAP.

## Target Variable:

- **Realized Volatility:**  
Square root of the sum of squared log returns, computed in **20-second buckets**.

# Data Split for Model Training

- Each time series (per `time_id`) was split **chronologically** into:
  - **First 80%** → used for training
  - **Last 20%** → held out for validation
- Ensures models are evaluated on **unseen future data**, simulating real-world forecasting.
- Prevents information leakage by maintaining the **temporal order** of volatility data.

# Candidate Models

1. Linear Regression

2. HAV-RV

3. eXtreme Gradient Boosting (XGBoost  
Tree-Based)

# Model 1 – Linear Regression

Call:

```
lm(formula = volatility ~ price + order + BidAskSpread, data = list.reg[[i]],  
    weights = 0.8^(((len.train - 2):0)/2))
```

## Objective:

Predict **realized volatility** using:

- **WAP (Weighted Average Price)** – proxy for execution price
- **Order Size** – total liquidity
- **Bid-Ask Spread** – market friction indicator

Uses **exponential weighting** to emphasize recent time buckets.

Built one model per time\_id to capture local patterns.

## Performance Summary (time\_id 162):

- **Adjusted R<sup>2</sup>: 0.497**

# Model 2 – HAV-RV Model

- Designed specifically for **high-frequency financial data**, the HAV-RV model forecasts volatility using past volatility patterns.
- Predictors:
  - vol\_1: volatility at time  $t-1$
  - mean\_vol\_5: average volatility over past 5 buckets
- Equation:

$$\hat{v}_t = \beta_0 + \beta_1 \cdot v_{t-1} + \beta_2 \cdot \bar{v}_{t-5:t-1}$$

- Also apply **Weighted Least Squares (WLS)** using quarticity to adjust for heteroskedasticity – i.e., varying levels of volatility uncertainty.
- At *time\_id* 162:
  - a. **WLS Adjusted R<sup>2</sup>: 0.4002**
  - b. **Unweighted Adjusted R<sup>2</sup>: 0.0293**

```
Call:  
lm(formula = vol ~ vol_1 + mean_vol_5, data = list.HAV[[i]])
```

```
Call:  
lm(formula = vol ~ vol_1 + mean_vol_5, data = list.HAV[[i]],  
    weights = list.HAV[[i]]$vol_1/sqrt(quar[[i]]$quarticity[5:(len.train -  
1)]))
```

# Model 3 – XGBoost Model: Non-Linear Tree-Based Regression

- Tree-based regression model
- Predicts volatility using market microstructure features
- Captures nonlinear patterns, outliers, and interactions
- No assumptions of stationarity/normality

## Feature Engineering

- Computed  $\log\_return = \log(WAP / \text{lag}(WAP))$  to measure price movement.
- Aggregated features per time\_id:
  - WAP\_mean, WAP\_sd
  - BidAskSpread\_mean, BidAskSpread\_sd
  - log\_return\_sd

## Target Variable

- Realized volatility =  $\sqrt{\text{sum}(\log\_return^2)}$

# Model Evaluation

## Evaluation Framework:



- **Data Split:** 80% training, 20% testing on time\_id-level features.
- **Validation:** First 100 time\_ids used for out-of-sample predictions for Linear, HAV-RV, and ARMA-GARCH. XGBoost evaluated on test set.
- **Metrics:**
  - **MAE / RMSE:** Measure average and squared error.
  - **R<sup>2</sup>:** Captures goodness-of-fit.
  - **QLIKE:** Financially motivated loss for volatility accuracy.



# Performance Results & Selection

	MAE	RMSE	R2	QLIKE
Linear	0.000315	0.000435	0.622229	0.168665
HAV_RV	0.000312	0.000444	0.601385	0.191400
XGBoost	0.000463	0.000728	0.953548	0.008318

# Model Results & Selection



# Final Model Selection: XGBoost

## Why XGBoost was selected:

### Best overall performance

- $R^2 = 0.95$ , QLIKE = 0.0083

### Strengths:

- Captures non-linear patterns, no distributional assumptions
- Robust to outliers and noise
- `log_return_sd` = most important feature

### Limitations

- It can overfit on noisy features if not regularized well
- Less interpretable than linear models

## Model Design:

- 80/20 split by `time_id`
- Engineered per `time_id`:
  - WAP, BidAskSpread, `log_return` (mean + sd)
- Target:
  - Realized volatility =  $\sqrt{(\sum \log\_return^2)}$
- Trained using:
  - 100 boosting rounds
  - `reg:squarederror` objective
  - Feature selection via `xgb.importance()`

Thank You For  
Listening  
Any Questions?